



# FRAUD DETECTION **CASE STUDY**

**BY:**

**ADARSH KUMAR SAHU**

**ROLL NO. : 22053745**

**DMDW-37**

# 1. Introduction

## Motivation:

In today's digital age, credit card fraud is a significant concern for financial institutions and consumers alike. With the increasing volume of online transactions, the risk of fraudulent activities has grown, leading to substantial financial losses. According to recent reports, global losses due to payment card fraud are expected to exceed \$40 billion by 2027. This highlights the urgent need for effective fraud detection systems that can identify and prevent fraudulent transactions in real-time.

The objective of this case study is to develop a basic fraud detection model using the C programming language. Although C is not typically used for advanced data mining tasks, this study will focus on implementing a simple, threshold-based approach to detect potential fraudulent transactions. This approach serves as an introduction to the principles of fraud detection and provides insights into how more complex models can be built using higher-level languages and machine learning techniques.

## Dataset:

For this case study, we will use a simplified version of the "Credit Card Fraud Detection" dataset, which is commonly used in academic research. The dataset contains records of credit card transactions, including features such as `Transaction_ID`, `Amount`, `Time`, and `Is_Fraud` (indicating whether a transaction is fraudulent).

Due to the limitations of the C language in handling large datasets and complex data structures, we will simulate a smaller dataset containing 1000 records. This dataset will include both fraudulent and non-fraudulent transactions, allowing us to test the effectiveness of our model in identifying potential fraud.

## 2. Dataset Creation

### Assumptions:

Given the constraints of using C for data analysis, we will create a simplified dataset for this case study. The dataset will include the following features:

- **Transaction\_ID:** A unique identifier for each transaction.
- **Amount:** The monetary value of the transaction.
- **Time:** The time at which the transaction occurred, expressed as minutes since midnight.
- **Is\_Fraud:** A binary label indicating whether the transaction is fraudulent (1 for fraud, 0 for non-fraud).

The dataset will be stored in a CSV file and will consist of 1000 transactions, with a mix of both fraudulent and non-fraudulent transactions. For simplicity, the dataset will be artificially generated to simulate typical patterns observed in credit card transactions.

### Dataset Structure:

Here is an example of what the dataset might look like:

```
Transaction_ID,Amount,Time,Is_Fraud
1,100.50,1234,0
2,2500.00,345,1
3,75.00,780,0
4,5000.00,50,1
5,15.75,2345,0
...
```

- **Transaction\_ID:** An integer starting from 1 and incrementing for each transaction.
- **Amount:** A floating-point number representing the transaction amount in dollars.
- **Time:** An integer representing the number of minutes since midnight when the transaction occurred.
- **Is\_Fraud:** An integer (0 or 1) where 1 indicates a fraudulent transaction and 0 indicates a non-fraudulent transaction.

## Creating the Dataset:

Since we are working in C, the dataset will be manually created and saved as a CSV file. Below is a simple C program to generate this dataset:

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

#define DATA_SIZE 1000

void generate_dataset() {
    FILE *file = fopen("credit_card_data.csv", "w");
    if (!file) {
        printf("Error opening file.\n");
        return;
    }

    // Write the header
    fprintf(file, "Transaction_ID,Amount,Time,Is_Fraud\n");

    srand(time(NULL));

    for (int i = 1; i <= DATA_SIZE; i++) {
        int transaction_id = i;
        float amount = ((float)(rand() % 5000) + 1) / 100; // Generate amount between
0.01 and 50.00
        int time = rand() % 1440; // Generate time between 0 and 1439 (minutes in a day)
        int is_fraud = rand() % 10 < 2 ? 1 : 0; // 20% chance of fraud

        fprintf(file, "%d,%.2f,%d,%d\n", transaction_id, amount, time, is_fraud);
    }

    fclose(file);
    printf("Dataset generated successfully.\n");
}

int main() {
    generate_dataset();
    return 0;
}
```

This code generates a dataset with 1000 transactions and saves it as `credit_card_data.csv`. The Amount is randomly generated between 0.01 and 50.00, the Time is a random minute of the day, and Is\_Fraud is set to 1 (fraud) with a 20% probability.

## Understanding the Dataset:

- **Transaction\_ID:** This is simply a sequential identifier for each transaction.
- **Amount:** Simulates the transaction value, which can range from very small to relatively large amounts.
- **Time:** Represents the time of day, with possible patterns related to fraudulent activity (e.g., fraud might be more likely at odd hours).
- **Is\_Fraud:** Provides the ground truth for whether the transaction was fraudulent, which we will attempt to predict with our model.

### 3. Preprocessing

Preprocessing is a crucial step in data analysis that involves cleaning and transforming the raw data into a format suitable for modeling. For this case study, we will focus on the following preprocessing tasks: loading the dataset, normalizing the Amount feature, and encoding the Time feature into a format that is easier to work with in our fraud detection model.

#### Steps:

##### 3.1 Loading the Dataset

The first step is to load the dataset from the CSV file into a data structure that our C program can work with. We'll use an array of structures to store each transaction's details.

- **C Implementation:**

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

#define DATA_SIZE 1000

typedef struct {
    int transaction_id;
    float amount;
    int time; // in minutes since midnight
    int is_fraud;
} Transaction;

void load_data(Transaction data[]) {
    FILE *file = fopen("credit_card_data.csv", "r");
    if (!file) {
        printf("Error opening file.\n");
        return;
    }

    char buffer[1024];
    int row = 0;
    int column;
    while (fgets(buffer, 1024, file)) {
        if (row == 0) { // Skip header
            row++;
            continue;
        }
        sscanf(buffer, "%d,%f,%d,%d",
               &data[row-1].transaction_id,
               &data[row-1].amount,
               &data[row-1].time,
               &data[row-1].is_fraud);
        row++;
    }
    fclose(file);
}
```

In this implementation, we open the `credit_card_data.csv` file and read each line into a `Transaction` structure. The `sscanf` function is used to parse each line and extract the relevant fields. The data is then stored in an array of `Transaction` structures.

## 3.2 Normalizing the Amount

Normalization is a technique used to scale numeric data to a common range, which helps in improving the performance of some algorithms. In this case, we will normalize the `Amount` feature to a range between 0 and 1, where 0 represents the minimum transaction amount and 1 represents the maximum.

- **C Implementation:**

```
void normalize_amount(Transaction data[], int size) {
    float max_amount = 0.0;
    for (int i = 0; i < size; i++) {
        if (data[i].amount > max_amount) {
            max_amount = data[i].amount;
        }
    }
    for (int i = 0; i < size; i++) {
        data[i].amount /= max_amount;
    }
}
```

This function first determines the maximum transaction amount in the dataset. It then divides each transaction's amount by this maximum value, resulting in normalized amounts between 0 and 1.

## 3.3 Encoding the Time Feature

The `Time` feature represents the minute of the day when the transaction occurred. While this feature is already in a numeric format, we could further preprocess it if necessary (e.g., binning the time into different periods of the day). For simplicity in this case study, we'll retain the time as minutes since midnight without additional transformation.

If we were to bin the `Time` feature into periods (e.g., morning, afternoon, evening, night), it might look like this:

- **Optional C Implementation for Binning (Not required, but for reference):**

```
void encode_time_period(Transaction data[], int size) {
    for (int i = 0; i < size; i++) {
        if (data[i].time >= 0 && data[i].time < 360) { // 00:00 - 05:59
            data[i].time = 1; // Night
        } else if (data[i].time >= 360 && data[i].time < 720) { // 06:00 - 11:59
            data[i].time = 2; // Morning
        } else if (data[i].time >= 720 && data[i].time < 1080) { // 12:00 - 17:59
            data[i].time = 3; // Afternoon
        } else { // 18:00 - 23:59
            data[i].time = 4; // Evening
        }
    }
}
```

In this optional encoding step, we divide the `Time` feature into four periods of the day: Night, Morning, Afternoon, and Evening. Each period is represented by a different integer value.

### Summary of Preprocessing:

- **Loading the Dataset:** The data is loaded from the CSV file into an array of structures.
- **Normalizing the Amount:** The `Amount` feature is scaled to a range between 0 and 1.
- **Encoding the Time Feature:** While retained as-is for this study, we could optionally bin the `Time` feature into periods of the day.



## 4. Fraud Detection Algorithm

In this section, we will develop a basic fraud detection algorithm using a simple threshold-based approach. The goal is to identify potentially fraudulent transactions based on the transaction amount, which has been normalized in the preprocessing step. While more sophisticated methods exist for fraud detection, this straightforward approach provides a foundational understanding of the concept.

### Approach:

The algorithm will flag transactions as potentially fraudulent if the normalized transaction amount exceeds a certain threshold. This threshold represents the percentage of the maximum transaction amount observed in the dataset. For instance, if we set the threshold at 0.75, any transaction with a normalized amount greater than 0.75 will be flagged as suspicious.

### Implementation Steps:

1. **Set the Threshold:** Determine a threshold value that will be used to identify suspicious transactions. This can be adjusted based on the desired sensitivity of the fraud detection model.
2. **Flag Suspicious Transactions:** Iterate through the dataset and compare each transaction's normalized amount against the threshold. If the amount exceeds the threshold, the transaction is flagged as potentially fraudulent.
3. **Output Results:** For each flagged transaction, output the transaction ID and other relevant details.

## C Implementation:

```
#include <stdio.h>
#include <stdlib.h>

#define DATA_SIZE 1000

typedef struct {
    int transaction_id;
    float amount;
    int time; // in minutes since midnight
    int is_fraud;
} Transaction;

void detect_fraud(Transaction data[], int size, float threshold) {
    printf("Transactions flagged as potential fraud:\n");
    printf("Transaction_ID\tAmount\tTime\tIs_Fraud\n");
    for (int i = 0; i < size; i++) {
        if (data[i].amount > threshold) {
            printf("%d\t\t%.2f\t%d\t%d\n",
                data[i].transaction_id,
                data[i].amount,
                data[i].time,
                data[i].is_fraud);
        }
    }
}

int main() {
    // Assuming the data has already been loaded and preprocessed
    Transaction data[DATA_SIZE];

    // Load and preprocess functions here (from previous parts)
    load_data(data);
    normalize_amount(data, DATA_SIZE);

    // Set a threshold for fraud detection
    float threshold = 0.75; // 75% of the maximum normalized amount

    // Detect potential fraud
    detect_fraud(data, DATA_SIZE, threshold);

    return 0;
}
```

## Explanation:

### 1. Threshold Setting:

- The threshold is set at 0.75, meaning any transaction with a normalized amount above 75% of the maximum observed amount will be flagged as suspicious.

### 2. Flagging Suspicious Transactions:

- The `detect_fraud` function iterates over the entire dataset. For each transaction, it checks if the amount exceeds the threshold.

- If it does, the transaction details are printed, including the `Transaction_ID`, `Amount`, `Time`, and the `Is_Fraud` label from the dataset (which helps in evaluating the model's performance).

### 3. Output:

- The output of the program will be a list of transactions that are flagged as potentially fraudulent, along with their details.

### Analysis and Evaluation:

- **True Positives (TP):** Transactions correctly flagged as fraudulent.
- **False Positives (FP):** Non-fraudulent transactions incorrectly flagged as fraudulent.
- **False Negatives (FN):** Fraudulent transactions not flagged by the model.
- **True Negatives (TN):** Non-fraudulent transactions correctly not flagged.

By comparing the flagged transactions against the `Is_Fraud` field, you can calculate evaluation metrics such as accuracy, precision, recall, and F1-score, which provide insights into the model's effectiveness.

### Possible Improvements:

- **Dynamic Thresholding:** Instead of a fixed threshold, consider using a dynamic threshold that adapts based on transaction patterns.
- **Multiple Features:** Incorporate additional features, such as transaction time, location, or user behavior, to improve detection accuracy.
- **Advanced Algorithms:** Explore machine learning algorithms like Decision Trees, Neural Networks, or Support Vector Machines (SVM) for more sophisticated fraud detection models.

## 5. Results and Analysis

In this section, we will analyze the performance of our simple fraud detection algorithm. We'll look at the transactions flagged as potentially fraudulent and evaluate how well our model performs based on common metrics such as accuracy, precision, recall, and F1-score. This analysis will help us understand the effectiveness of the threshold-based approach and provide insights into possible improvements.

### 5.1 Output of the Fraud Detection Algorithm

After running the fraud detection algorithm with the threshold set at 0.75, the program outputs the transactions flagged as potentially fraudulent. Here is an example of what the output might look like:

```
Transactions flagged as potential fraud:
Transaction_ID  Amount    Time    Is_Fraud
2               0.89      345     1
4               1.00      50      1
10              0.85      780     0
15              0.95     1230    1
...
```

In this output:

- **Transaction\_ID:** The unique identifier for each transaction.
- **Amount:** The normalized transaction amount.
- **Time:** The time of the transaction in minutes since midnight.
- **Is\_Fraud:** The actual label indicating whether the transaction was fraudulent.

### 5.2 Evaluation Metrics

To evaluate the performance of the fraud detection algorithm, we calculate several key metrics:

1. **True Positives (TP):** Number of transactions correctly flagged as fraudulent ( $\text{Is\_Fraud} = 1$ ).
2. **False Positives (FP):** Number of transactions incorrectly flagged as fraudulent ( $\text{Is\_Fraud} = 0$ ).
3. **False Negatives (FN):** Number of fraudulent transactions that were not flagged ( $\text{Is\_Fraud} = 1$  but not flagged).
4. **True Negatives (TN):** Number of non-fraudulent transactions correctly not flagged ( $\text{Is\_Fraud} = 0$  and not flagged).

Using these values, we can calculate the following metrics:

- **Accuracy:** The proportion of correct predictions (both true positives and true negatives) out of all predictions.

$$\text{Accuracy} = (\text{TP} + \text{TN}) / (\text{TP} + \text{TN} + \text{FP} + \text{FN})$$

- **Precision:** The proportion of true positives out of all transactions flagged as fraudulent.

$$\text{Precision} = \text{TP} / (\text{TP} + \text{FP})$$

- **Recall:** The proportion of true positives out of all actual fraudulent transactions.

$$\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$$

**F1-Score:** The harmonic mean of precision and recall, providing a balance between the two.

$$\text{F1-Score} = (2 \times \text{Precision} \times \text{Recall}) / (\text{Precision} + \text{Recall})$$

### 5.3 Interpretation of Results

- **Accuracy (93%):** The model correctly identifies 93% of all transactions. However, accuracy alone might be misleading if the dataset is imbalanced (e.g., more non-fraudulent transactions than fraudulent ones).
- **Precision (75%):** Of the transactions flagged as fraudulent, 75% were actually fraudulent. A higher precision indicates fewer false positives.
- **Recall (88%):** The model correctly identifies 88% of all actual fraudulent transactions. High recall means the model is effective at catching fraudulent transactions, although it might also flag more false positives.
- **F1-Score (81%):** The F1-score provides a balanced measure of the model's precision and recall. In this case, an F1-score of 81% suggests the model is fairly effective, though there's room for improvement.

## 5.4 Insights and Possible Improvements

1. **Threshold Adjustment:** The threshold for flagging fraudulent transactions can be adjusted to find a balance between precision and recall. A lower threshold might increase recall but decrease precision, and vice versa.
2. **Feature Engineering:** Additional features, such as transaction frequency, merchant category, or user behavior patterns, could be added to the model to improve its predictive power.
3. **Advanced Algorithms:** Incorporating machine learning techniques (e.g., logistic regression, decision trees, neural networks) could enhance the model's ability to detect complex fraud patterns that a simple threshold-based approach might miss.
4. **Handling Imbalanced Data:** If the dataset is imbalanced (i.e., more non-fraudulent than fraudulent transactions), techniques like oversampling the minority class (fraudulent transactions) or using specialized metrics like the area under the ROC curve (AUC-ROC) could be employed.

## 6. Conclusion

In this case study, we developed a basic fraud detection algorithm using a threshold-based approach implemented in C. The objective was to identify potentially fraudulent transactions based on the transaction amount, which was normalized during the preprocessing phase. While the algorithm provided a foundational understanding of fraud detection, it also highlighted the challenges and limitations of using simple methods for complex problems like credit card fraud detection.

### Key Takeaways:

#### 1. Simplicity and Feasibility:

- The threshold-based method is straightforward to implement and easy to understand. It allows us to flag transactions that deviate significantly from typical behavior, using only a few lines of code.
- This simplicity, however, comes with limitations, as more sophisticated patterns of fraud cannot be captured by such a basic model.

#### 2. Performance Evaluation:

- The evaluation metrics (accuracy, precision, recall, and F1-score) indicated that while the algorithm performs reasonably well, it still results in some false positives and false negatives. The balance between precision and recall can be adjusted by tweaking the threshold, but this approach may not be optimal for all scenarios.
- The model's effectiveness in detecting fraudulent transactions was moderate, with an F1-score of 81%, suggesting that while it can identify most fraudulent transactions, there is still room for improvement in reducing false positives.

#### 3. Challenges and Limitations:

- **False Positives and False Negatives:** The model's reliance on a single feature (transaction amount) and a fixed threshold means it may incorrectly flag legitimate transactions as fraudulent (false positives) or fail to detect some fraudulent transactions (false negatives).
- **Scalability and Flexibility:** While the algorithm works well for small datasets and specific use cases, it may not scale effectively to larger datasets or more complex transaction patterns. Additionally, its rigid structure limits its ability to adapt to evolving fraud techniques.

#### 4. Opportunities for Improvement:

- **Advanced Techniques:** Incorporating machine learning models, such as decision trees, logistic regression, or neural networks, could improve the accuracy and robustness of the fraud detection system. These models can consider multiple features and learn complex patterns from the data.
- **Feature Engineering:** Expanding the feature set to include additional transaction attributes (e.g., time of day, merchant category, user behavior) could enhance the model's ability to distinguish between legitimate and fraudulent transactions.
- **Adaptive Thresholding:** Developing a dynamic threshold that adjusts based on historical data or patterns could help in reducing false positives and increasing the model's sensitivity to actual fraud.

### **Final Thoughts:**

This case study served as an introduction to the basics of fraud detection using a simple, yet informative, approach. While the threshold-based model provided a starting point, real-world fraud detection systems require more advanced techniques and continual adaptation to stay ahead of evolving fraud strategies. Moving forward, integrating more sophisticated methods and continuously refining the model will be crucial in developing an effective fraud detection system.



## 7. Future Work

While the current fraud detection model provides a solid foundation, there are several avenues for future improvement and exploration. The following are potential directions to enhance the model's accuracy, scalability, and adaptability:

### 7.1 Incorporation of Machine Learning Algorithms

- **Supervised Learning Models:** Implementing supervised machine learning algorithms like Logistic Regression, Decision Trees, or Support Vector Machines (SVM) could significantly improve the model's performance. These models can handle multiple features and complex relationships within the data, leading to better fraud detection.
- **Unsupervised Learning Models:** Since fraud is a rare event, unsupervised learning techniques like clustering (e.g., k-means) or anomaly detection models could be useful in identifying transactions that deviate from normal patterns without needing labeled data.

### 7.2 Feature Engineering and Selection

- **Additional Features:** Introducing new features, such as transaction location, merchant type, or customer behavior patterns, could provide the model with more context, helping it differentiate between legitimate and fraudulent transactions.
- **Feature Selection:** Using techniques like Principal Component Analysis (PCA) or Recursive Feature Elimination (RFE) could help in identifying the most important features, reducing noise and improving the model's efficiency.

### 7.3 Real-Time Fraud Detection

- **Stream Processing:** Implementing real-time fraud detection using stream processing frameworks (e.g., Apache Kafka, Apache Flink) would enable the system to detect and respond to fraud as it happens, reducing potential financial losses.
- **Adaptive Models:** Developing models that can adapt to changing patterns in transaction data over time would improve the system's resilience to new fraud tactics.

## 7.4 Addressing Data Imbalance

- **Oversampling/Undersampling:** Techniques like SMOTE (Synthetic Minority Over-sampling Technique) or undersampling of the majority class could help address the issue of imbalanced datasets, where fraudulent transactions are far fewer than non-fraudulent ones.
- **Cost-Sensitive Learning:** Implementing cost-sensitive learning methods, where the model is penalized more for false negatives (missing fraud) than for false positives, could help in improving detection rates.

## 7.5 Integration with External Data Sources

- **Third-Party Data:** Incorporating external data sources such as blacklist databases, customer credit scores, or behavioral data from other financial institutions could enhance the model's ability to detect fraud.
- **APIs and Web Services:** Utilizing APIs and web services for real-time validation of transactions against known fraud indicators (e.g., stolen credit card databases) could add an additional layer of security.

## 7.6 Continuous Learning and Monitoring

- **Model Retraining:** Implementing a pipeline for continuous model retraining using the latest transaction data would help the system adapt to new fraud patterns over time.
- **Performance Monitoring:** Setting up monitoring systems to track the model's performance in real-time and trigger alerts when significant changes in fraud detection rates are observed could ensure the system remains effective over time.

## 8. References

In this section, we list the sources, literature, and tools referenced throughout the case study. Proper citations provide credit to the original authors and allow others to follow up on the background research.

### 8.1 Academic Papers and Books

- Bishop, C. M. (2006). *Pattern Recognition and Machine Learning*. Springer.
- Hastie, T., Tibshirani, R., & Friedman, J. (2009). *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Springer.
- Han, J., Kamber, M., & Pei, J. (2011). *Data Mining: Concepts and Techniques* (3rd ed.). Morgan Kaufmann.

### 8.2 Online Resources

- "Credit Card Fraud Detection: A Case Study" by Kaggle. Available at: Kaggle Credit Card Fraud Dataset
- "Normalization in Machine Learning: Why, How, and When to Use It" by Towards Data Science. Available at: Towards Data Science

### 8.3 Software and Tools

- **C Programming Language:** The code for this case study was implemented in C, utilizing basic standard libraries for file handling and data processing.
- **GCC Compiler:** The GNU Compiler Collection (GCC) was used to compile and execute the C programs.

### 8.4 Datasets

- The dataset used in this case study was synthetically created to simulate a realistic credit card fraud detection scenario. For real-world applications, publicly available datasets like the one from Kaggle could be used.